

# AmazonFineFoodReviewsAnalysisDecisionTree

February 11, 2019

## 1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

**Objective:** Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

## 2 [1]. Reading Data

### 2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [145]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import pickle as pkl
import os
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc, roc_auc_score
from sklearn.tree import DecisionTreeClassifier

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors

from tqdm import tqdm
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import TimeSeriesSplit, GridSearchCV

from prettytable import PrettyTable
from IPython.display import Image

In [4]: # Read the Amazon fine food review data from database using sqlite
con = sqlite3.connect('database.sqlite')

# Select all reviews where score is not 3 (neutral)
review_data = pd.read_sql_query("""SELECT * FROM Reviews WHERE Score != 3""", con)
```

```
# Assign positive class if score >=4 else assign negative class
score = review_data['Score']
PN_score = score.map(lambda x: "Positive" if x>=4 else "Negative")
review_data['Score'] = PN_score

print("Shape of review data is {}".format(review_data.shape))
review_data.head(3)
```

Shape of review data is (525814, 10)

```
Out[4]:
```

	Id	ProductId	UserId	ProfileName	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	
2	3	B000LQOCHO	ABXLMWJIXXAIN	Natalia Corres	"Natalia Corres"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	1	1	Positive	1303862400	
1	0	0	Negative	1346976000	
2	1	1	Positive	1219017600	

	Summary	Text
0	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	"Delight" says it all	This is a confection that has been around a fe...

```
In [5]: #Trying to visualize the duplicate data before removal
display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [7]: print(display.shape)
display.head()
```

(80668, 7)

```
Out[7]:
```

	UserId	ProductId	ProfileName	Time	Score	\
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory	1342396800	5	"hoppy"
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	
3	#oc-R1105J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	

	Text	COUNT(*)
--	------	----------

0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [8]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out[8]:
```

	UserId	ProductId	ProfileName	Time \
	80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine" 1334707200

	Score	Text	COUNT(*)
80638	5	I was recommended to try green tea extract to ...	5

```
In [9]: display['COUNT(*)'].sum()
```

```
Out[9]: 393063
```

## 3 [2] Exploratory Data Analysis

### 3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [12]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out[12]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator \
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600
1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

Summary \

```

0  LOACKER QUADRATINI VANILLA WAFERS
1  LOACKER QUADRATINI VANILLA WAFERS
2  LOACKER QUADRATINI VANILLA WAFERS
3  LOACKER QUADRATINI VANILLA WAFERS
4  LOACKER QUADRATINI VANILLA WAFERS

```

```

                                Text
0  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

```

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```

In [8]: #Sorting data according to ProductId in ascending order
        #sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)

        #Deduplication of entries
        #final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
        #final.shape

```

```

In [13]: #Remove the duplicate entries from the data

        sorted_data = review_data.sort_values('ProductId')
        final = sorted_data.drop_duplicates(subset=["UserId", "Time", "Summary"])
        print(final.shape)

```

```

(363186, 10)

```

```

In [15]: #Checking to see how much % of data still remains
        print((final['Id'].size*1.0)/(review_data['Id'].size*1.0)*100)

```

```

69.07119247490557

```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [16]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

```
Out[16]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	3	1	5	1224892800	
1	3	2	4	1212883200	

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [18]: # Removing the reviews where HelpfulnessNumerator > HelpfulnessDenominator
```

```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [20]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
print(final['Score'].value_counts())
```

```
(363184, 10)
Positive    306173
Negative    57011
Name: Score, dtype: int64
```

## 4 [3] Preprocessing

### 4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [42]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
# https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup
```

```
In [47]: # https://stackoverflow.com/a/47091490/4084039
import re
```

```
def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\'re", " are", phrase)
    phrase = re.sub(r"\'s", " is", phrase)
    phrase = re.sub(r"\'d", " would", phrase)
    phrase = re.sub(r"\'ll", " will", phrase)
    phrase = re.sub(r"\'t", " not", phrase)
    phrase = re.sub(r"\'ve", " have", phrase)
    phrase = re.sub(r"\'m", " am", phrase)
    return phrase
```

```
In [43]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
#remove spacial character: https://stackoverflow.com/a/5843547/4084039
```

```
In [116]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
```

*# instead of <br /> if we have <br/> these tags would have removed in the 1st step*

```
stopwords = set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourse',
                "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
                'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
                'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that',
                'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has',
                'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'a',
                'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
                'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off',
                'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'a',
                'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 't',
                's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've",
                've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
                'hadn't', 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'm',
                'mustn't', 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
                'won', "won't", 'wouldn', "wouldn't"])
```

In [55]: `from nltk.stem import SnowballStemmer`

```
#Intializing SnowballStemmer
snow_stemmer = SnowballStemmer('english')

#Using Stemmer on a word
print(snow_stemmer.stem('Moves'))
```

move

In [48]: *# Combining all the above to clean reviews*

```
from tqdm import tqdm
preprocessed_reviews = []

# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

100%|| 363184/363184 [04:45<00:00, 1270.29it/s]

In [64]: *# Storing the preprocessed reviews and stemmed preprocessed reviews seperately.*  
*# We have performed the cleaning on the whole data so we can use it later on*



```

# models other than KNN that can handle high dimensional data gracefully.

#####---- storing the data into .sqlite file -----#####
# Reviews are present in preprocessed_reviews

final['CleanedText'] = preprocessed_reviews

#Store the data into a sqlite database
if not os.path.isfile('final.sqlite'):
    conn = sqlite3.connect('final.sqlite')
    c = conn.cursor()
    conn.text_factory = str
    final.to_sql('Reviews', conn, schema=None, if_exists='replace', \
                index=True, index_label=None, chunksize=None, dtype=None)
    conn.close()

In [91]: # Performing stemming on the preprocessed reviews
final['CleanedText'] = preprocessed_reviews
stemmed_reviews = []

for sentence in final['CleanedText'].values:
    sentence = b' '.join((snow_stemmer.stem(word)).encode('utf8') for word in sentence)
    stemmed_reviews.append(sentence)

In [85]: final['CleanedText'] = stemmed_reviews
final['CleanedText'] = final['CleanedText'].str.decode("utf-8")

if not os.path.isfile('final_stemmedreviews.sqlite'):
    conn = sqlite3.connect('final_stemmedreviews.sqlite')
    c = conn.cursor()
    conn.text_factory = str
    final.to_sql('Reviews', conn, schema=None, if_exists='replace', \
                index=True, index_label=None, chunksize=None, dtype=None)
    conn.close()

In [3]: # Load the preprocessed dataset from the database final.sqlite
# Data is ordered by time stamps to facilitate time base splitting
# of data for cross validation

conn = sqlite3.connect('final.sqlite')
final = pd.read_sql_query("""SELECT * From Reviews ORDER BY Time""", conn)
conn.close()

conn = sqlite3.connect('final_stemmedreviews.sqlite')
final_stemmed = pd.read_sql_query("""SELECT * From Reviews ORDER BY Time""", conn)
conn.close()

In [4]: # There is an extra index column in the data
final.head(1)

```

```

Out[4]:      index      Id  ProductId      UserId      ProfileName \
0  138706  150524  0006641040  ACITT7DI6IDDL  shari zychinski

      HelpfulnessNumerator  HelpfulnessDenominator      Score      Time \
0                        0                        0  Positive  939340800

      Summary \
0  EVERY book is educational

      Text \
0  this witty little book makes my son laugh at l...

      CleanedText
0  witty little book makes son laugh loud recite ...

```

```

In [5]: #Removing the index column from data
clean_data = final.drop(['index'], axis=1)

#clean_data_stemmed = final_stemmed.drop(['index'], axis=1)

# Map postive to 1 and negative to 0 in Score column
score = clean_data['Score']
bin_score = score.map(lambda x: 1 if x == "Positive" else 0)
clean_data['Score'] = bin_score

# Add stemmed reviews as an extra column in the data
# This will be in addition to the preprocessed non stemmed
# reviews which are stored in the CleanedText column.

stemmed_reviews = final_stemmed['CleanedText']
clean_data['StemmedText'] = stemmed_reviews

```

```

In [6]: # Adding another feature into the data
# we will find the length of the each review
# and add that as a feature into the existing
# dataframe.

clean_data['Reviewlen'] = clean_data['StemmedText'].apply(len)

```

```

In [7]: clean_data.tail(1)

```

```

Out[7]:      Id  ProductId      UserId ProfileName  HelpfulnessNumerator \
363183  5703  B009WSNWC4  AMP7K1084DH1T      ESTY                        0

      HelpfulnessDenominator  Score      Time      Summary \
363183                        0      1  1351209600  DELICIOUS

      Text \
363183  Purchased this product at a local store in NY ...

```

	CleanedText \	
363183	purchased product local store ny kids love qui...	
	StemmedText	Reviewlen
363183	purchas product local store ny kid love quick ...	109

```
In [8]: # Split the dataset in training and test dataset
# We will use the training data for cross validation and training.
# Test data will not be known to model and will be used
# to calculate the accuracy.
```

```
# Data is split in 70-30 train-test split using slicing since
# data is sorted in ascending time order
```

```
# Instead of splitting the data and then sampling
# let's try to split the 100k samples directly and
# then just simple time split the data in 70-30k
```

```
data = clean_data.iloc[:,:]
subset_data = data.iloc[100000:200000,:]
```

```
train_cv_split = 70000
```

```
train = subset_data.iloc[:train_cv_split,:]
test = subset_data.iloc[train_cv_split:,:]
```

```
print(train.shape , '\n', test.shape)
```

```
(70000, 13)
(30000, 13)
```

```
In [9]: print(train[train['Score'] == 0].shape)
print(test[test['Score'] == 0].shape)
```

```
(11235, 13)
(4961, 13)
```

```
In [10]: # Seperating the Score column from rest of the data
columns = list(clean_data.columns)
columns = [column for column in columns if column != 'Score']

X_train = train[columns]
y_train = train['Score']

X_test = test[columns]
y_test = test['Score']
```

```

print(X_train.shape , y_train.shape, '\n', X_test.shape, y_test.shape)

(70000, 12) (70000,)
(30000, 12) (30000,)

```

```

In [11]: # Save the y_train and y_test so we
         # can directly use it later rather than rerunning
         # the splitting steps again

pk1.dump(y_train, open("y_train.pk1", 'wb'))
pk1.dump(y_test, open("y_test.pk1", 'wb'))

```

### [3.2] Preprocessing Review Summary

```

In [6]: ## Similarly you can do preprocessing for review summary also.

```

## 5 [4] Featurization

### 5.1 [4.1] BAG OF WORDS

```

In [ ]: # Obtaining a vectorizer on stemmed reviews
         # It was observed during Word2Vec transformation
         # that stemmed reviews give words which are close to
         # say good or bad otherwise we observe other words
         # which seem non-relevant. So we will use stemmed reviews.

# Words in stemmed review that are most similar to great and worst
#[('wonder', 0.7626501321792603), ('awesom', 0.7493463754653931), ('excel', 0.74750399
#('fantast', 0.7294141054153442), ('good', 0.7276639938354492), ('terrif', 0.696876645
#('nice', 0.6279305219650269), ('perfect', 0.6089357733726501), ('amaz', 0.57377290725
#('decent', 0.5731742978096008)]
#=====
#[('horribl', 0.7659773826599121), ('disgust', 0.7506155967712402), ('terribl', 0.7292
#('aw', 0.7216229438781738), ('nasti', 0.6849608421325684), ('foul', 0.661132156848907
#('gag', 0.6592600345611572), ('weird', 0.6567815542221069), ('funni', 0.6493463516235
#('gross', 0.6418379545211792)]

# Words in stemmed review that are most similar to great and worst
# As we can see worst is similar to greatest and best in non-stemmed reviews.

#[('awesome', 0.7547115087509155), ('fantastic', 0.7433849573135376), ('wonderful', 0.
#('excellent', 0.7240736484527588), ('good', 0.7088381052017212), ('terrific', 0.66505
#('amazing', 0.6410914659500122), ('perfect', 0.6294776201248169), ('fabulous', 0.6247
#('incredible', 0.5898726582527161)]
#=====
#[('greatest', 0.7661513090133667), ('best', 0.668804407119751), ('richest', 0.6509857

```

```

#('smoothest', 0.6451543569564819), ('nastiest', 0.639174222946167), ('tastiest', 0.61
#('encountered', 0.6121875047683716), ('disgusting', 0.600991427898407), ('yummiest',
#('nicest', 0.5876485705375671)]

```

```

In [14]: # Running count vectorizer on training data only
# to avoid data leakage
# we will use the uni-grams & bi-grams in BoW embedding
# count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
count_vec = CountVectorizer(ngram_range=(1,2), min_df=10)

X_train_bow = count_vec.fit_transform(X_train['StemmedText'].values)
X_test_bow = count_vec.transform(X_test['StemmedText'].values)

# Save the training and test BOW vectors in pickle files
# We can simply load this data later and use it

pkl.dump(X_train_bow, open("train_bow.pkl", 'wb'))
pkl.dump(X_test_bow, open("test_bow.pkl", 'wb'))
pkl.dump(count_vec, open("count_vec.pkl", 'wb'))

```

## 5.2 [4.2] TF-IDF

```

In [16]: # Apply tfidf vectorizer to convert text to vectors

tf_idf = TfidfVectorizer(ngram_range=(1,2), min_df=10)

X_train_tfidf = tf_idf.fit_transform(X_train['StemmedText'].values)
X_test_tfidf = tf_idf.transform(X_test['StemmedText'].values)

# Save the training, CV and test TFIDF vectors in pickle files
# We can simply load this data later and use it

pkl.dump(X_train_tfidf, open("train_tfidf.pkl", 'wb'))
pkl.dump(X_test_tfidf, open("test_tfidf.pkl", 'wb'))
pkl.dump(tf_idf, open("tf_idf.pkl", 'wb'))

In [17]: # Creating a dictionary with word as key and it's tfidf representation as value
dictionary = dict(zip(tf_idf.get_feature_names(), list(tf_idf.idf_)))

pkl.dump(dictionary, open("tfidf_dictionary.pkl", 'wb'))

```

## 5.3 [4.3] Word2Vec

```

In [11]: # Train our own Word2Vec model using your own text corpus

list_of_sent_test = []
list_of_sent_train = []

```

```

for review in X_test['StemmedText'].values:
    list_of_sent_test.append(review.split())

for review in X_train['StemmedText'].values:
    list_of_sent_train.append(review.split())

w2v = Word2Vec(list_of_sent_train, min_count=5, size=100, workers=4)
w2v.save('w2v_model.bin')
w2v_words = list(w2v.wv.vocab)

In [12]: print(w2v.wv.most_similar('great'))
print('='*50)
print(w2v.wv.most_similar('bad'))

[('fantast', 0.7587853670120239), ('excel', 0.7455682158470154), ('wonder', 0.7229946255683899),
=====
[('horribl', 0.706422746181488), ('terribl', 0.7024113535881042), ('aw', 0.674425482749939), (

In [13]: w2v_words = list(w2v.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:100])

number of words that occurred minimum 5 times 11131
sample words ['hey', 'good', 'stuff', 'like', 'tasti', 'cold', 'hot', 'flavor', 'subtl', 'yet

```

## 5.4 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

### [4.4.1.1] Avg W2v

```

In [117]: # Avg-W2V
from tqdm import tqdm

train_review_vectors = []
test_review_vectors = []

dataset = [(list_of_sent_train, train_review_vectors),
           (list_of_sent_test, test_review_vectors)]

for item in dataset:
    for review in tqdm(item[0]):
        nwords = 0
        rev_vec = np.zeros(100)
        for word in review:
            if word in w2v_words:
                vec = w2v.wv[word]
                rev_vec += vec
                nwords += 1

```

```

        if nwords != 0:
            rev_vec /= nwords
            item[1].append(rev_vec)

```

```

100%|| 70000/70000 [01:41<00:00, 686.62it/s]
100%|| 30000/30000 [00:43<00:00, 686.55it/s]

```

In [118]: *# Save the review vectors so we can use later*

```

pk1.dump(train_review_vectors, open("train_avgw2v.pkl", 'wb'))
pk1.dump(test_review_vectors, open("test_avgw2v.pkl", 'wb'))

```

#### [4.4.1.2] TFIDF weighted W2v

In [14]: `tf_idf = TfidfVectorizer(ngram_range=(1,2), min_df=10)`

```

X_train_tfidf = tf_idf.fit_transform(X_train['StemmedText'].values)
X_test_tfidf = tf_idf.transform(X_test['StemmedText'].values)

dictionary_tfidf = dict(zip(tf_idf.get_feature_names(), list(tf_idf.idf_)))
tfidf_features = tf_idf.get_feature_names()

```

In [15]: *# review\_vectors will store the tfidf-weighted W2V representation of the reviews in t*

```

# TFIDFWeighted-W2V
from tqdm import tqdm

train_review_vectors = []
test_review_vectors = []

list_of_sent_test = []
list_of_sent_train = []

for review in X_test['CleanedText'].values:
    list_of_sent_test.append(review.split())

for review in X_train['CleanedText'].values:
    list_of_sent_train.append(review.split())

dataset = [(list_of_sent_train, train_review_vectors),
            (list_of_sent_test, test_review_vectors)]

w2v_model = Word2Vec.load('w2v_model.bin')
w2v_words = list(w2v_model.wv.vocab)

for item in dataset:

```

```

row=0
for review in tqdm(item[0]):
    rev_vec = np.zeros(100)
    weight_sum = 0
    for word in review:
        if word in w2v_words and word in tfidf_features:
            vec = w2v_model.wv[word]
            tf_idf = dictionary_tfidf[word]*(review.count(word)/len(review))
            rev_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        rev_vec /= weight_sum
    item[1].append(rev_vec)
    row += 1

```

100%|| 70000/70000 [40:56<00:00, 28.50it/s]

100%|| 30000/30000 [16:10<00:00, 30.92it/s]

In [16]: *# Save the review vectors so we can use later*

```

pkl.dump(train_review_vectors, open("train_tfidfw2v.pkl", 'wb'))
pkl.dump(test_review_vectors, open("test_tfidfw2v.pkl", 'wb'))

```

## 5.5 Utility Functions used in Decision Trees

In [95]: *# This function takes the vector representation of review data  
# and returns the optimal depth of the tree and minimum samples  
# left till we split for Decision Tree using 5-fold  
# cross validation in GridSearchCV.  
# Code below makes use of TimeSeriesSplit.*

```

def get_optimal_hyperparams(X_train, y_train):
    parameters = {'max_depth' : [1, 5, 10, 50, 100, 500, 1000], \
                  'min_samples_split' : [5, 10, 100, 500]}

    #Perform GridSearch
    cv_obj = TimeSeriesSplit(n_splits=5).split(X_train)
    clf = GridSearchCV(DecisionTreeClassifier(), parameters,
                       scoring = 'roc_auc', cv=cv_obj)
    clf.fit(X_train, y_train)

    #tree_depth = parameters['max_depth']
    #min_samples = parameters['min_samples_split']
    gresults = clf.cv_results_
    hyper_params = gresults['params']

    auc_scores_train = gresults['mean_train_score']

```



```

auc_scores_cv = gresults['mean_test_score']
#print("cv_results : {}".format(gresults))
#print("Best : {}".format(clf.best_score_))
#optimal_c = c_values[results['rank_test_score'][0]-1]
optimal_depth = gresults['params'][clf.best_index_]['max_depth']
optimal_minsplit = gresults['params'][clf.best_index_]['min_samples_split']
#print(clf.best_estimator_)
#print(clf.best_index_, '\t', optimal_depth, '\t', optimal_minsplit)
return optimal_depth, optimal_minsplit, zip(hyper_params,
                                            auc_scores_train, auc_scores_cv)

```

In [96]: *# Running Decision Tree Classifier with given max\_depth  
# and min\_samples\_split value and returns a tuple indicating  
# AUC obtained for test data along with the confusion matrix  
# along with the classifier object. Same function can be used  
# on all vectorized data irrespective of vectorizer*

```

def run_dt(X_train, y_train, X_test, y_test, depth, minsplit):
    dt_clf = DecisionTreeClassifier(max_depth=depth,
                                    min_samples_split= minsplit)

    dt_clf.fit(X_train, y_train)

    y_pred_test = dt_clf.predict_proba(X_test)
    y_pred_train = dt_clf.predict_proba(X_train)

    y_pred_test_prob = y_pred_test[:,1]
    y_pred_test_label = np.argmax(y_pred_test, axis=1)

    y_pred_train_prob = y_pred_train[:,1]
    y_pred_train_label = np.argmax(y_pred_train, axis=1)

    fpr_test, tpr_test, thresholds_test = roc_curve(y_test,\
                                                    y_pred_test_prob)
    auc_score_test = auc(fpr_test, tpr_test)

    fpr_train, tpr_train, thresholds_train = roc_curve(y_train, \
                                                    y_pred_train_prob)
    auc_score_train = auc(fpr_train, tpr_train)

    conf_mat = confusion_matrix(y_test, y_pred_test_label)

    plt.figure()
    plt.plot(fpr_train, tpr_train, color='darkorange', lw=1, \
             label='Train ROC curve (area = %0.2f)' % auc_score_train)
    plt.plot(fpr_test, tpr_test, color='navy', lw=1, \
             label='Test ROC curve (area = %0.2f)' % auc_score_test)
    plt.plot([0, 1], [0, 1], color='black', lw=1, linestyle='--')
    plt.xlim([0.0, 1.0])

```

```

plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.show()
plt.close()
return dt_clf, auc_score_test, conf_mat

In [88]: def plot_confusion_matrix(cm):
    labels = ['Negative', 'Positive']
    confmat = pd.DataFrame(cm, index = labels, columns = labels)
    sns.heatmap(confmat, annot = True, fmt = 'd', cmap="Greens")
    plt.title("Confusion Matrix")
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.show()

In [86]: # All the results will be stored in the results dataframe and
# later in we will use this dataframe to print the results
# in tabular format

results = pd.DataFrame(
    columns=['Features-Used', 'Vectorizer', 'Max Depth', 'Min Samples Split', 'AUC'])

In [87]: import prettytable as pt
# function to print the results obtained in a table format
def print_results(data):
    result = PrettyTable(hrules=pt.ALL,
                        vrules=pt.ALL, padding_width=2)
    result.field_names = list(data.columns)
    for i in range(0, data.shape[0]):
        result.add_row(data.iloc[i])
    print(result)

In [74]: # Function to plot the heatmap of AUC scores
# for validation and training auc scores against
# the parameter grid i.e. max_depth and min_samples_split
def plot_heatmap(params, auc_scores, flag):
    max_depth = sorted(set([item['max_depth'] \
                            for item in params]))
    min_samples_split = sorted(set([item['min_samples_split'] \
                                    for item in params]))

    auc_scores = np.array(auc_scores).reshape(len(min_samples_split),
                                              len(max_depth))

    auc_scores_mat = pd.DataFrame(auc_scores,
                                  index=min_samples_split,
                                  columns=max_depth)

```

```

clr = sns.color_palette("Blues")
ax = sns.heatmap(auc_scores_mat, annot = True, cmap=clr)
ax.invert_yaxis()
plt.title("Heat Map for {} AUC scores".format(flag))
plt.xlabel("Max Depth")
plt.ylabel("Minimum Sample Split")
plt.show()

```

## 6 Applying Decision Trees

### 6.1 [5.1] Applying Decision Trees on BOW, SET 1

```

In [97]: # Load the saved vectorized data for train-test datapoints
X_train_bow = pickle.load(open('train_bow.pkl', 'rb'))
X_test_bow = pickle.load(open('test_bow.pkl', 'rb'))

# Getting an optimal value of hyperparameter max_depth
# and min_samples_split. This data is used to plot a heatmap
# There will be a problem of data leakage while using
# Gridsearch on train data but no way to get around it.
# Test data doesn't have this problem since it is transformed
# using the vectorizer fit on training data.
depth, minsplit, auc_scores = get_optimal_hyperparams(X_train_bow,
                                                        y_train)

params, train_auc, cv_auc = zip(*auc_scores)
print("Optimal Max Depth : {}".format(depth))
print("Optimal Min Samples Split : {}".format(minsplit))

# Plotting hyperparameter values vs AUC scores
plot_heatmap(params, train_auc, flag="Train")
plot_heatmap(params, cv_auc, flag="CV")

# Running Decision Tree Classifier with optimal value
# of max_depth and min_samples_split obtained
dt_clf, auc_score, conf_mat = run_dt(X_train_bow, y_train,
                                       X_test_bow, y_test,
                                       depth, minsplit)

print("AUC score:\n {:.2f}".format(auc_score))

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)


auc_score = '%0.2f' % auc_score



```

```
# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "BoW", depth, minsplit, auc_score]
```


Optimal Max Depth : 50

Optimal Min Samples Split : 500





AUC score:  
0.85

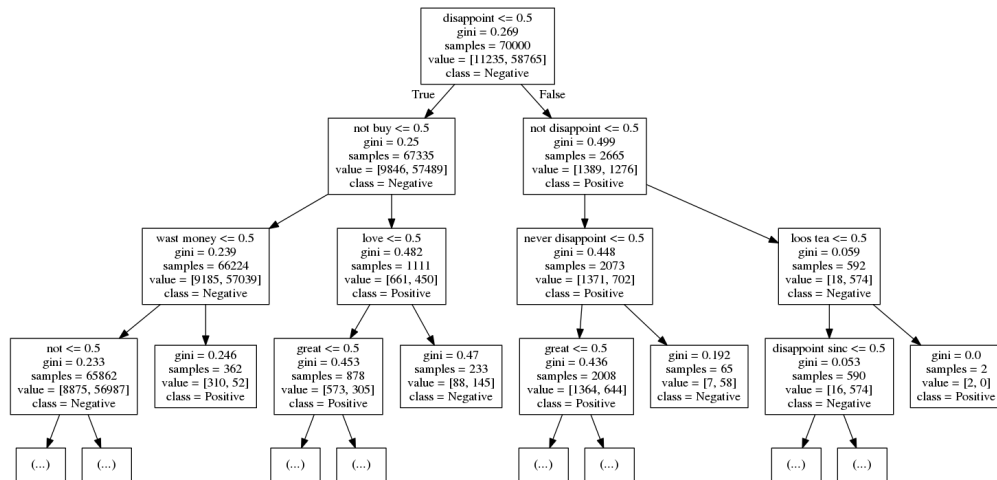


### 6.1.1 [5.1.1] Top 20 important features from SET 1

```
In [129]: # Loading the saved count vectorizer object
count_vec = pickle.load(open("count_vec.pkl", 'rb'))

# Top 20 features as per feature importance
top_indices = np.argsort(-dt_clf.feature_importances_)[:20, None]
print(np.take(count_vec.get_feature_names(), top_indices))
```

```
[['disappoint']
 ['not buy']
 ['wast money']
 ['not disappoint']
 ['great']
 ['not']
 ['love']
 ['worst']
 ['return']
 ['best']
 ['aw']]
```



title

```

['horribl']
['delici']
['refund']
['good']
['not worth']
['not good']
['not recommend']
['terribl']
['perfect']]

```

## 6.1.2 [5.1.2] Graphviz visualization of Decision Tree on BOW, SET 1

In [132]: *# The Graphical representation of decision tree  
# is being saved in a png file named tree\_bow.png  
# Restricting the tree to depth 3.*

```

from sklearn.tree import export_graphviz

export_graphviz(dt_clf, out_file='dt_graph_bow', max_depth=3,
                feature_names=count_vec.get_feature_names(),
                class_names=['Negative', 'Positive'])

```

## 6.1.3 Graphical representation of Decision Tree(BoW)

## 6.2 [5.2] Applying Decision Trees on TFIDF, SET 2

In [139]: *# Load the saved vectorized data for train-test datapoints*  
 X\_train\_tfidf = pkl.load(open('train\_tfidf.pkl', 'rb'))  
 X\_test\_tfidf = pkl.load(open('test\_tfidf.pkl', 'rb'))

```

# Getting an optimal value of hyperparameter max_depth
# and min_samples_split. This data is used to plot a heatmap
# There will be a problem of data leakage while using
# Gridsearch on train data but no way to get around it.
# Test data doesn't have this problem since it is transformed
# using the vectorizer fit on training data.
depth, minsplit, auc_scores = get_optimal_hyperparams(X_train_tfidf,
                                                    y_train)

params, train_auc, cv_auc = zip(*auc_scores)
print("Optimal Max Depth : {}".format(depth))
print("Optimal Min Samples Split : {}".format(minsplit))

# Plotting hyperparameter values vs AUC scores
plot_heatmap(params, train_auc, flag="Train")
plot_heatmap(params, cv_auc, flag="CV")

# Running Decision Tree Classifier with optimal value
# of max depth and min samples split obtained
dt_clf, auc_score, conf_mat = run_dt(X_train_tfidf, y_train,
                                     X_test_tfidf, y_test,
                                     depth, minsplit)

print("AUC score:\n {:.2f}".format(auc_score))

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)

auc_score = '%0.2f' % auc_score

# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "TF-IDF", depth, minsplit, auc_score]


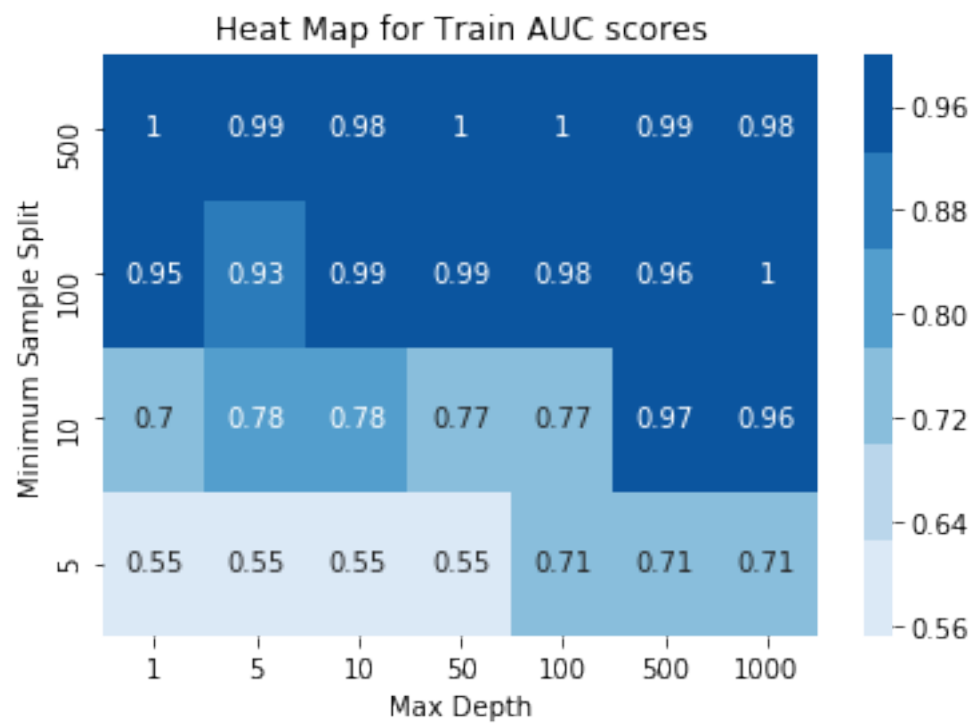
```


```

Optimal Max Depth : 50
Optimal Min Samples Split : 500


```







AUC score:  
0.84



### 6.2.1 [5.1.1] Top 20 important features from SET 2

```
In [140]: # Loading the saved tf-idf vectorizer object
          tf_idf = pickle.load(open("tf_idf.pkl", 'rb'))

          # Top 20 features as per feature importance
          top_indices = np.argsort(-dt_clf.feature_importances_)[:20, None]
          print(np.take(tf_idf.get_feature_names(), top_indices))

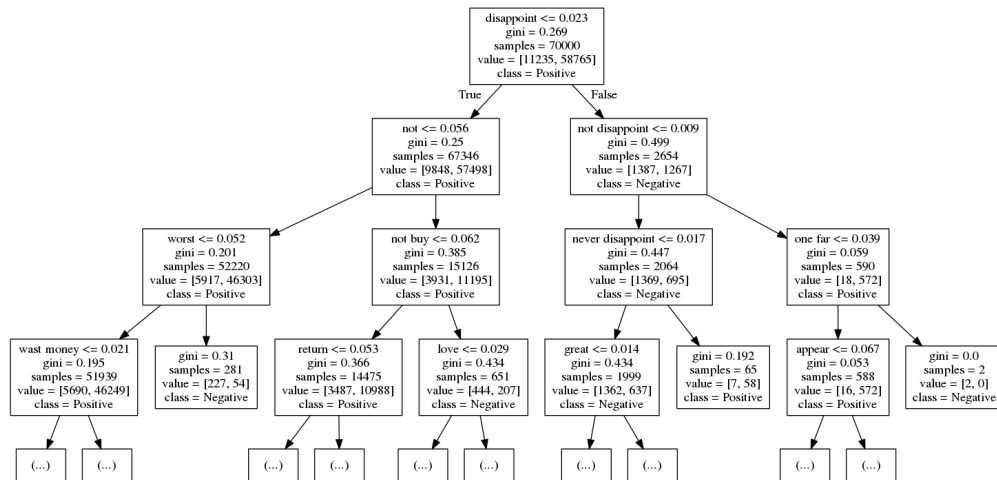
[['disappoint']
 ['not']
 ['not disappoint']
 ['great']
 ['not buy']
 ['worst']
 ['return']
 ['love']
 ['wast money']
 ['aw']
 ['horribl']
 ['refund']
 ['best']
 ['threw']
 ['good']
 ['terribl']
 ['delici']
 ['money']
 ['disgust']
 ['not worth']]
```

### 6.2.2 [5.2.2] Graphviz visualization of Decision Tree on TFIDF, SET 2

```
In [142]: # The Graphical representation of decision tree
          # is being saved in a png file named tree_tfidf.png
          # Restricting the tree to depth 3.

          from sklearn.tree import export_graphviz

          export_graphviz(dt_clf, out_file='dt_graph_tfidf', max_depth=3,
                          feature_names=count_vec.get_feature_names(),
                          class_names=['Negative', 'Positive'])
```



title

### 6.2.3 Graphical representation of Decision Tree(TF-IDF)

## 6.3 [5.3] Applying Decision Trees on AVG W2V SET 3

```
In [135]: # Load the saved vectorized data for train-test datapoints
X_train_avgw2v = pickle.load(open('train_avgw2v.pkl', 'rb'))
X_test_avgw2v = pickle.load(open('test_avgw2v.pkl', 'rb'))

# Getting an optimal value of hyperparameter max_depth
# and min_samples_split. This data is used to plot a heatmap
# There will be a problem of data leakage while using
# Gridsearch on train data but no way to get around it.
# Test data doesn't have this problem since it is transformed
# using the vectorizer fit on training data.
depth, minsplit, auc_scores = get_optimal_hyperparams(X_train_avgw2v,
                                                        y_train)

params, train_auc, cv_auc = zip(*auc_scores)
print("Optimal Max Depth : {}".format(depth))
print("Optimal Min Samples Split : {}".format(minsplit))

# Plotting hyperparameter values vs AUC scores
plot_heatmap(params, train_auc, flag="Train")
plot_heatmap(params, cv_auc, flag="CV")

# Running Decision Tree Classifier with optimal value
# of max depth and min samples split obtained
dt_clf, auc_score, conf_mat = run_dt(X_train_avgw2v, y_train,
                                       X_test_avgw2v, y_test,
                                       depth, minsplit)

print("AUC score:\n {:.2f}".format(auc_score))
```

```

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)


auc_score = '%0.2f' % auc_score



# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "AvgW2V", depth, minsplit, auc_score]

```


Optimal Max Depth : 10

Optimal Min Samples Split : 500





AUC score:  
0.82



## 6.4 [5.4] Applying Decision Trees on TFIDF W2V SET 4

```
In [136]: # Load the saved vectorized data for train-test datapoints
X_train_tfidf2v = pickle.load(open('train_tfidf2v.pkl', 'rb'))
X_test_tfidf2v = pickle.load(open('test_tfidf2v.pkl', 'rb'))

# Getting an optimal value of hyperparameter max_depth
# and min_samples_split. This data is used to plot a heatmap.
# There will be a problem of data leakage while using
# Gridsearch on train data but no way to get around it.
# Test data doesn't have this problem since it is transformed
# using the vectorizer fit on training data.
depth, minsplit, auc_scores = get_optimal_hyperparams(X_train_tfidf2v,
                                                       y_train)

params, train_auc, cv_auc = zip(*auc_scores)
print("Optimal Max Depth : {}".format(depth))
print("Optimal Min Samples Split : {}".format(minsplit))

# Plotting hyperparameter values vs AUC scores
```

```

plot_heatmap(params, train_auc, flag="Train")
plot_heatmap(params, cv_auc, flag="CV")

# Running Decision Tree Classifier with optimal value
# of max depth and min samples split obtained
dt_clf, auc_score, conf_mat = run_dt(X_train_tfidf2v, y_train,
                                     X_test_tfidf2v, y_test,
                                     depth, minsplit)

print("AUC score:\n {:.2f}".format(auc_score))

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)


auc_score = '%0.2f' % auc_score

# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "Tfidf-W2V", depth, minsplit, auc_score]



```

Optimal Max Depth : 10


Optimal Min Samples Split : 500







AUC score:  
0.75



## 7 [6] Conclusions

1. We tried BoW, TF-IDF, Average Word2Vec and Tfidf weighted Word2Vec vectorizers on Decision Tree Classifier.
2. AUC score for Decision Tree classifier was better for BoW and TFIDF vectorizers.

3. BoW :

Top 20 important features

'disappoint' , 'not buy', 'wast money', 'not disappoint', 'great', 'not',  
'love', 'worst', 'return', 'best', 'aw', 'horribl', 'delici', 'refund',  
'good', 'not worth', 'not good', 'not recommend', 'terribl', 'perfect'

4. TFIDF :

Top 20 important features

'disappoint', 'not', 'not disappoint', 'great', 'not buy', 'worst', 'return',  
'love', 'wast money', 'aw', 'horribl', 'refund', 'best', 'threw', 'good',  
'terribl', 'delici', 'money', 'disgust', 'not worth'

```
In [152]: print_results(results)
```

Features-Used	Vectorizer	Max Depth	Min Samples Split	AUC
Review Text	BoW	50	500	0.85
Review Text	TF-IDF	50	500	0.84
Review Text	AvgW2V	10	500	0.82
Review Text	Tfidf-W2V	10	500	0.75