

AmazonFineFoodReviewsAnalysisLogRegression

February 5, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [3]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import pickle as pkl
import os
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc, roc_auc_score
from sklearn.linear_model import LogisticRegression

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors

from tqdm import tqdm
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import TimeSeriesSplit, GridSearchCV

from prettytable import PrettyTable

In [4]: # Read the Amazon fine food review data from database using sqlite
con = sqlite3.connect('database.sqlite')

# Select all reviews where score is not 3 (neutral)
review_data = pd.read_sql_query("""SELECT * FROM Reviews WHERE Score != 3""", con)

# Assign positive class if score >=4 else assign negative class
```

```

score = review_data['Score']
PN_score = score.map(lambda x: "Positive" if x>=4 else "Negative")
review_data['Score'] = PN_score

print("Shape of review data is {}".format(review_data.shape))
review_data.head(3)

```

Shape of review data is (525814, 10)

```

Out[4]:
   Id  ProductId  UserId  ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK      dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      1                      1  Positive  1303862400
1                      0                      0  Negative  1346976000
2                      1                      1  Positive  1219017600

           Summary  Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1    Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [5]: #Trying to visualize the duplicate data before removal
display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [7]: print(display.shape)
display.head()

```

(80668, 7)

```

Out[7]:
   UserId  ProductId  ProfileName  Time  Score \
0  #oc-R115TNMSPFT9I7  B007Y59HVM      Breyton  1331510400      2
1  #oc-R11D9D7SHXIJB9  B005HG9ET0  Louis E. Emory "hoppy"  1342396800      5
2  #oc-R11DNU2NBKQ23Z  B007Y59HVM      Kim Cieszykowski  1348531200      1
3  #oc-R1105J5ZVQE25C  B005HG9ET0      Penguin Chick  1346889600      5
4  #oc-R12KPBODL2B5ZD  B0070SBE1U  Christopher P. Presta  1348617600      1

           Text  COUNT(*)
0  Overall its just OK when considering the price...      2

```

1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [8]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [8]:
```

	UserId	ProductId	ProfileName	Time \
	80638	AZY10LLTJ71NX	B006P7E5ZI undertheshrine	"undertheshrine" 1334707200

	Score	Text	COUNT(*)
80638	5	I was recommended to try green tea extract to ...	5

```
In [9]: display['COUNT(*)'].sum()
```

```
Out [9]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [12]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out [12]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator \
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600
1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

```
Summary \
0 LOCKER QUADRATINI VANILLA WAFERS
```

```

1  LOACKER QUADRATINI VANILLA WAFERS
2  LOACKER QUADRATINI VANILLA WAFERS
3  LOACKER QUADRATINI VANILLA WAFERS
4  LOACKER QUADRATINI VANILLA WAFERS

```

```

                                Text
0  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

```

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```

In [8]: #Sorting data according to ProductId in ascending order
        #sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)

        #Deduplication of entries
        #final=sorted_data.drop_duplicates(subset={"UserId", "ProfileName", "Time", "Text"}, keep='first')
        #final.shape

```

```

In [13]: #Remove the duplicate entries from the data

        sorted_data = review_data.sort_values('ProductId')
        final = sorted_data.drop_duplicates(subset=["UserId", "Time", "Summary"])
        print(final.shape)

(363186, 10)

```

```

In [15]: #Checking to see how much % of data still remains
        print((final['Id'].size*1.0)/(review_data['Id'].size*1.0)*100)

69.07119247490557

```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [16]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)
```

```
display.head()
```

```
Out[16]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	3	1	5	1224892800	
1	3	2	4	1212883200	

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [18]: # Removing the reviews where HelpfullnessNumerator > HelpfulnessDenominator
```

```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [20]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
print(final['Score'].value_counts())
```

```
(363184, 10)
Positive    306173
Negative    57011
Name: Score, dtype: int64
```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [42]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
# https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup
```

```
In [47]: # https://stackoverflow.com/a/47091490/4084039
import re
```

```
def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase
```

```
In [43]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
#remove spacial character: https://stackoverflow.com/a/5843547/4084039
```

```
In [116]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have reumoved in the 1st step

stopwords = set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourse',
'you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that',
'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has',
'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'a
```

```
'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off',
'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'each',
'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've",
've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
'hadn't', 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'me',
'mustn't', 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
'won', "won't", 'wouldn', "wouldn't"])
```

```
In [55]: from nltk.stem import SnowballStemmer
```

```
#Intializing SnowballStemmer
snow_stemmer = SnowballStemmer('english')
```

```
#Using Stemmer on a word
print(snow_stemmer.stem('Moves'))
```

move

```
In [48]: # Combining all the above to clean reviews
```

```
from tqdm import tqdm
preprocessed_reviews = []
```

```
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

```
100%|| 363184/363184 [04:45<00:00, 1270.29it/s]
```

```
In [64]: # Storing the preprocessed reviews and stemmed preprocessed reviews seperately.
# We have performed the cleaning on the whole data so we can use it later on
# models other than KNN that can handle high dimensional data gracefully.
```

```
#####---- storing the data into .sqlite file -----#####
# Reviews are present in preprocessed_reviews
```

```
final['CleanedText'] = preprocessed_reviews
```

```
#Store the data into a sqlite database
```



```

if not os.path.isfile('final.sqlite'):
    conn = sqlite3.connect('final.sqlite')
    c = conn.cursor()
    conn.text_factory = str
    final.to_sql('Reviews', conn, schema=None, if_exists='replace', \
                index=True, index_label=None, chunksize=None, dtype=None)
    conn.close()

```

In [91]: *# Performing stemming on the preprocessed reviews*

```

final['CleanedText'] = preprocessed_reviews
stemmed_reviews = []

```

```

for sentence in final['CleanedText'].values:
    sentence = b' '.join((snow_stemmer.stem(word)).encode('utf8') for word in sentence)
    stemmed_reviews.append(sentence)

```

In [85]: *final['CleanedText'] = stemmed_reviews*

```

final['CleanedText'] = final['CleanedText'].str.decode("utf-8")

```

```

if not os.path.isfile('final_stemmedreviews.sqlite'):
    conn = sqlite3.connect('final_stemmedreviews.sqlite')
    c = conn.cursor()
    conn.text_factory = str
    final.to_sql('Reviews', conn, schema=None, if_exists='replace', \
                index=True, index_label=None, chunksize=None, dtype=None)
    conn.close()

```

In [4]: *# Load the preprocessed dataset from the database final.sqlite*

Data is ordered by time stamps to facilitate time base splitting
of data for cross validation

```

conn = sqlite3.connect('final.sqlite')
final = pd.read_sql_query("""SELECT * From Reviews ORDER BY Time""", conn)
conn.close()

```

```

conn = sqlite3.connect('final_stemmedreviews.sqlite')
final_stemmed = pd.read_sql_query("""SELECT * From Reviews ORDER BY Time""", conn)
conn.close()

```

In [5]: *# There is an extra index column in the data*

```

final.head(1)

```

Out[5]:

	index	Id	ProductId	UserId	ProfileName	\
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	0	0	Positive	939340800	

Summary \

```
0 EVERY book is educational
```

```
Text \
0 this witty little book makes my son laugh at 1...
```

```
CleanedText
0 witty little book makes son laugh loud recite ...
```

```
In [6]: #Removing the index column from data
clean_data = final.drop(['index'], axis=1)

#clean_data_stemmed = final_stemmed.drop(['index'], axis=1)

# Map postive to 1 and negative to 0 in Score column
score = clean_data['Score']
bin_score = score.map(lambda x: 1 if x == "Positive" else 0)
clean_data['Score'] = bin_score

# Add stemmed reviews as an extra column in the data
# This will be in addition to the preprocessed non stemmed
# reviews which are stored in the CleanedText column.

stemmed_reviews = final_stemmed['CleanedText']
clean_data['StemmedText'] = stemmed_reviews
```

```
In [7]: # Adding another feature into the data
# we will find the length of the each review
# and add that as a feature into the existing
# dataframe.

clean_data['Reviewlen'] = clean_data['StemmedText'].apply(len)
```

```
In [8]: clean_data.tail(1)
```

```
Out[8]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	
363183	5703	B009WSNWC4	AMP7K1084DH1T	ESTY	0	
	HelpfulnessDenominator	Score	Time	Summary		
363183	0	1	1351209600	DELICIOUS		
				Text		
363183				Purchased this product at a local store in NY ...		
				CleanedText		
363183				purchased product local store ny kids love qui...		
				StemmedText	Reviewlen	
363183				purchas product local store ny kid love quick ...	109	

```

In [9]: # Split the dataset in training and test dataset
        # We will use the training data for cross validation and training.
        # Test data will not be known to model and will be used
        # to calculate the accuracy.

        # Data is split in 70-30 train-test split using slicing since
        # data is sorted in ascending time order

        # Instead of splitting the data and then sampling
        # let's try to split the 100k samples directly and
        # then just simple time split the data in 70-30k

data = clean_data.iloc[:, :]
subset_data = data.iloc[100000:200000, :]

train_cv_split = 70000

train = subset_data.iloc[:train_cv_split, :]
test = subset_data.iloc[train_cv_split:, :]

print(train.shape , '\n', test.shape)

(70000, 13)
(30000, 13)

In [10]: print(train[train['Score'] == 0].shape)
         print(test[test['Score'] == 0].shape)

(11235, 13)
(4961, 13)

In [11]: # Separating the Score column from rest of the data
columns = list(clean_data.columns)
columns = [column for column in columns if column != 'Score']

X_train = train[columns]
y_train = train['Score']

X_test = test[columns]
y_test = test['Score']

print(X_train.shape , y_train.shape, '\n', X_test.shape, y_test.shape)

(70000, 12) (70000,)
(30000, 12) (30000,)

```

```
In [12]: # Save the y_train and y_test so we
        # can directly use it later rather than rerunning
        # the splitting steps again

        pickle.dump(y_train, open("y_train.pkl", 'wb'))
        pickle.dump(y_test, open("y_test.pkl", 'wb'))
```

[3.2] Preprocessing Review Summary

```
In [6]: ## Similarly you can do preprocessing for review summary also.
```

5 [4] Featurization

5.1 [4.1] BAG OF WORDS

```
In [ ]: # Obtaining a vectorizer on stemmed reviews
        # It was observed during Word2Vec transformation
        # that stemmed reviews give words which are close to
        # say good or bad otherwise we observe other words
        # which seem non-relevant. So we will use stemmed reviews.

        # Words in stemmed review that are most similar to great and worst
        #[('wonder', 0.7626501321792603), ('awesom', 0.7493463754653931), ('excel', 0.74750399),
        #('fantast', 0.7294141054153442), ('good', 0.7276639938354492), ('terrifi', 0.696876645),
        #('nice', 0.6279305219650269), ('perfect', 0.6089357733726501), ('amaz', 0.57377290725),
        #('decent', 0.5731742978096008)]
        #=====
        #[('horribl', 0.7659773826599121), ('disgust', 0.7506155967712402), ('terribl', 0.7292),
        #('aw', 0.7216229438781738), ('nasti', 0.6849608421325684), ('foul', 0.661132156848907),
        #('gag', 0.6592600345611572), ('weird', 0.6567815542221069), ('funni', 0.6493463516235),
        #('gross', 0.6418379545211792)]

        # Words in stemmed review that are most similar to great and worst
        # As we can see worst is similar to greatest and best in non-stemmed reviews.

        #[('awesome', 0.7547115087509155), ('fantastic', 0.7433849573135376), ('wonderful', 0.7433849573135376),
        #('excellent', 0.7240736484527588), ('good', 0.7088381052017212), ('terrific', 0.66505),
        #('amazing', 0.6410914659500122), ('perfect', 0.6294776201248169), ('fabulous', 0.6247),
        #('incredible', 0.5898726582527161)]
        #=====
        #[('greatest', 0.7661513090133667), ('best', 0.668804407119751), ('richest', 0.6509857),
        #('smoothest', 0.6451543569564819), ('naughtiest', 0.639174222946167), ('tastiest', 0.61),
        #('encountered', 0.6121875047683716), ('disgusting', 0.600991427898407), ('yummy', 0.600991427898407),
        #('nicest', 0.5876485705375671)]

In [14]: # Running count vectorizer on training data only
        # to avoid data leakage
        # we will use the uni-grams & bi-grams in BoW embedding
```

```

# count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
count_vec = CountVectorizer(ngram_range=(1,2), min_df=10)

X_train_bow = count_vec.fit_transform(X_train['StemmedText'].values)
X_test_bow = count_vec.transform(X_test['StemmedText'].values)

# Save the training and test BOW vectors in pickle files
# We can simply load this data later and use it

pkl.dump(X_train_bow, open("train_bow.pkl", 'wb'))
pkl.dump(X_test_bow, open("test_bow.pkl", 'wb'))
pkl.dump(count_vec, open("count_vec.pkl", 'wb'))

```

5.2 [4.2] TF-IDF

In [16]: # Apply tfidf vectorizer to convert text to vectors

```

tf_idf = TfidfVectorizer(ngram_range=(1,2), min_df=10)

X_train_tfidf = tf_idf.fit_transform(X_train['StemmedText'].values)
X_test_tfidf = tf_idf.transform(X_test['StemmedText'].values)

# Save the training, CV and test TFIDF vectors in pickle files
# We can simply load this data later and use it

pkl.dump(X_train_tfidf, open("train_tfidf.pkl", 'wb'))
pkl.dump(X_test_tfidf, open("test_tfidf.pkl", 'wb'))
pkl.dump(tf_idf, open("tf_idf.pkl", 'wb'))

```

In [17]: # Creating a dictionary with word as key and it's tfidf representation as value

```

dictionary = dict(zip(tf_idf.get_feature_names(), list(tf_idf.idf_)))

```

```

pkl.dump(dictionary, open("tfidf_dictionary.pkl", 'wb'))

```

5.3 [4.3] Word2Vec

In [11]: # Train our own Word2Vec model using your own text corpus

```

list_of_sent_test = []
list_of_sent_train = []

for review in X_test['StemmedText'].values:
    list_of_sent_test.append(review.split())

for review in X_train['StemmedText'].values:
    list_of_sent_train.append(review.split())

w2v = Word2Vec(list_of_sent_train, min_count=5, size=100, workers=4)

```

```

w2v.save('w2v_model.bin')
w2v_words = list(w2v.wv.vocab)

In [12]: print(w2v.wv.most_similar('great'))
print('='*50)
print(w2v.wv.most_similar('bad'))

[('fantast', 0.7587853670120239), ('excel', 0.7455682158470154), ('wonder', 0.7229946255683899),
=====
[('horribl', 0.706422746181488), ('terribl', 0.7024113535881042), ('aw', 0.674425482749939), (

In [13]: w2v_words = list(w2v.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:100])

number of words that occurred minimum 5 times 11131
sample words ['hey', 'good', 'stuff', 'like', 'tasti', 'cold', 'hot', 'flavor', 'subtl', 'yet

```

5.4 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```

In [117]: # Avg-W2V
from tqdm import tqdm

train_review_vectors = []
test_review_vectors = []

dataset = [(list_of_sent_train, train_review_vectors),
           (list_of_sent_test, test_review_vectors)]

for item in dataset:
    for review in tqdm(item[0]):
        nwords = 0
        rev_vec = np.zeros(100)
        for word in review:
            if word in w2v_words:
                vec = w2v.wv[word]
                rev_vec += vec
                nwords += 1
        if nwords != 0:
            rev_vec /= nwords
        item[1].append(rev_vec)

100%|| 70000/70000 [01:41<00:00, 686.62it/s]
100%|| 30000/30000 [00:43<00:00, 686.55it/s]

```

```
In [118]: # Save the review vectors so we can use later
```

```
pk1.dump(train_review_vectors, open("train_avgw2v.pkl", 'wb'))
pk1.dump(test_review_vectors, open("test_avgw2v.pkl", 'wb'))
```

[4.4.1.2] TFIDF weighted W2v

```
In [14]: tf_idf = TfidfVectorizer(ngram_range=(1,2), min_df=10)
```

```
X_train_tfidf = tf_idf.fit_transform(X_train['StemmedText'].values)
X_test_tfidf = tf_idf.transform(X_test['StemmedText'].values)

dictionary_tfidf = dict(zip(tf_idf.get_feature_names(), list(tf_idf.idf_)))
tfidf_features = tf_idf.get_feature_names()
```

```
In [15]: # review_vectors will store the tfidf-weighted W2V representation of the reviews in t
```

```
# TFIDFWeighted-W2V
from tqdm import tqdm

train_review_vectors = []
test_review_vectors = []

list_of_sent_test = []
list_of_sent_train = []

for review in X_test['CleanedText'].values:
    list_of_sent_test.append(review.split())

for review in X_train['CleanedText'].values:
    list_of_sent_train.append(review.split())

dataset = [(list_of_sent_train, train_review_vectors),
            (list_of_sent_test, test_review_vectors)]

w2v_model = Word2Vec.load('w2v_model.bin')
w2v_words = list(w2v_model.wv.vocab)

for item in dataset:
    row=0
    for review in tqdm(item[0]):
        rev_vec = np.zeros(100)
        weight_sum = 0
        for word in review:
            if word in w2v_words and word in tfidf_features:
                vec = w2v_model.wv[word]
                tf_idf = dictionary_tfidf[word]*(review.count(word)/len(review))
```

```

        rev_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        rev_vec /= weight_sum
    item[1].append(rev_vec)
    row += 1

```

```

100%|| 70000/70000 [40:56<00:00, 28.50it/s]
100%|| 30000/30000 [16:10<00:00, 30.92it/s]

```

In [16]: *# Save the review vectors so we can use later*

```

pkl.dump(train_review_vectors, open("train_tfidfv2v.pkl", 'wb'))
pkl.dump(test_review_vectors, open("test_tfidfv2v.pkl", 'wb'))

```

5.5 Utility Functions used in Logistic Regression

In [13]: *# This function takes the vector representation of review data
and returns the optimal C for Logistic Regression using 5-fold
cross validation in GridSearchCV.
Code below makes use of TimeSeriesSplit.*

```

def get_optimal_c(X_train, y_train, penalty='l2'):
    parameters = {'C' : [10**i for i in range(-5,6)]}

    #Perform GridSearch
    cv_obj = TimeSeriesSplit(n_splits=5).split(X_train)
    clf = GridSearchCV(LogisticRegression(penalty=penalty), parameters,
                        scoring = 'roc_auc', cv=cv_obj)
    clf.fit(X_train, y_train)

    c_values = parameters['C']
    gresults = clf.cv_results_
    auc_scores_train = gresults['mean_train_score']
    auc_scores_cv = gresults['mean_test_score']
    #print("cv_results : {}".format(results))
    #print("Best : {}".format(clf.best_score_))
    #optimal_c = c_values[results['rank_test_score'][0]-1]
    optimal_c = gresults['params'][clf.best_index_]['C']
    return optimal_c, zip(c_values, auc_scores_train),\
        zip(c_values, auc_scores_cv)

```

In [14]: *# Running Logistic Regression with given C
returns a tuple indicating AUC obtained for
test data along with the confusion matrix along with the classifier
object same function can be used on all vectorized data irrespective
of vectorizer*


```

def run_lr(X_train, y_train, X_test, y_test, c, penalty='l2'):
    lr_clf = LogisticRegression(penalty=penalty,C=c)
    lr_clf.fit(X_train, y_train)

    y_pred_test = lr_clf.predict_proba(X_test)
    y_pred_train = lr_clf.predict_proba(X_train)

    y_pred_test_prob = y_pred_test[:,1]
    y_pred_test_label = np.argmax(y_pred_test, axis=1)

    y_pred_train_prob = y_pred_train[:,1]
    y_pred_train_label = np.argmax(y_pred_train, axis=1)

    fpr_test, tpr_test, thresholds_test = roc_curve(y_test,\
                                                    y_pred_test_prob)
    auc_score_test = auc(fpr_test, tpr_test)

    fpr_train, tpr_train, thresholds_train = roc_curve(y_train, \
                                                        y_pred_train_prob)
    auc_score_train = auc(fpr_train, tpr_train)

    conf_mat = confusion_matrix(y_test, y_pred_test_label)

    plt.figure()
    plt.plot(fpr_train, tpr_train, color='darkorange', lw=1, \
             label='Train ROC curve (area = %0.2f)' % auc_score_train)
    plt.plot(fpr_test, tpr_test, color='navy', lw=1, \
             label='Test ROC curve (area = %0.2f)' % auc_score_test)
    plt.plot([0, 1], [0, 1], color='black', lw=1, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.0])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver operating characteristic')
    plt.legend(loc="lower right")
    plt.show()
    plt.close()
    return lr_clf, auc_score_test, conf_mat

```

```

In [15]: def plot_confusion_matrix(cm):
    labels = ['Negative', 'Positive']
    confmat = pd.DataFrame(cm, index = labels, columns = labels)
    sns.heatmap(confmat, annot = True, fmt = 'd', cmap="Greens")
    plt.title("Confusion Matrix")
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.show()

```

```

In [16]: def get_sparsity(clf_obj):

```

```

        # Number of non zero elements in weight vector
        non_zero = np.count_nonzero(clf_obj.coef_)
        total = clf_obj.coef_.shape[1]
        sparsity = round((total - non_zero) / float(total), 4)
        return sparsity*100

In [17]: # All the results will be stored in the results dataframe and
        # later in we will use this dataframe to print the results
        # in tabular format

        results = pd.DataFrame(
            columns=['Features-Used', 'Vectorizer', 'Regularizer', 'C', 'Sparsity', 'AUC'])

In [26]: import prettytable as pt
        # function to print the results obtained in a table format
        def print_results(data):
            result = PrettyTable(hrules=pt.ALL,
                                vrules=pt.ALL, padding_width=2)
            result.field_names = list(data.columns)
            for i in range(0, data.shape[0]):
                result.add_row(data.iloc[i])
            print(result)

```

6 Applying Logistic Regression

6.1 [5.1] Logistic Regression on BOW, SET 1

6.1.1 [5.1.1] Applying Logistic Regression with L1 regularization on BOW, SET 1

```

In [19]: # Load the saved vectorized data for train-test datapoints
        X_train_bow = pickle.load(open('train_bow.pkl', 'rb'))
        X_test_bow = pickle.load(open('test_bow.pkl', 'rb'))

        std = StandardScaler(with_mean=False)

        # Standardizing the vectors
        X_train_bow_std = std.fit_transform(X_train_bow)
        X_test_bow_std = std.transform(X_test_bow)

        # Getting an optimal value of hyperparameter c and AUC scores
        # This data is used to plot a graph of C-values vs AUC
        # There will be a problem of data leakage while using
        # Gridsearch on train data but no way to get around it.
        # Test data doesn't have this problem since it is transformed
        # using the vectorizer fit on training data.
        optimal_c, c_auc_train, c_auc_cv = get_optimal_c(X_train_bow_std,
                                                         y_train,
                                                         penalty='l1')

```

```

print("Optimal value of C : {}".format(optimal_c))

train_auc = [(c, train_auc) for c, train_auc in c_auc_train]
cv_auc = [(c, cv_auc) for c, cv_auc in c_auc_cv]

# Plotting c values vs AUC scores

plt.title("C vs AUC")
plt.xlabel("Hyperparameter(c)")
plt.ylabel("Area under ROC curve")
plt.plot(*zip(*train_auc), label='Train_AUC')
plt.plot(*zip(*cv_auc), label='Validation_AUC')
plt.legend()
plt.show()

# Running logistic Regression with optimal alpha value obtained
lr_clf, auc_score, conf_mat = run_lr(X_train_bow_std, y_train,
                                     X_test_bow_std, y_test,
                                     optimal_c, penalty='l1')

print("AUC score:\n {:.2f}".format(auc_score))

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)



auc_score = '%0.2f' % auc_score

# Sparsity of weight vector to results
sparsity = get_sparsity(lr_clf)


# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "BoW", \
                                'l1', optimal_c, sparsity, auc_score]

```

Optimal value of C : 0.01



AUC score:
0.95



[5.1.1.1] Calculating sparsity on weight vector obtained using L1 regularization on BOW, SET 1

```
In [20]: print("Sparsity of weight vector : {}".format(sparsity))
```

Sparsity of weight vector : 91.38%

6.1.2 [5.1.2] Applying Logistic Regression with L2 regularization on BOW, SET 1

```
In [21]: # Changing the regularizer to l2
```

```
# Getting an optimal value of hyperparameter c and AUC scores
# This data is used to plot a graph of C-values vs AUC
# There will be a problem of data leakage while using
# Gridsearch on train data but no way to get around it.
# Test data doesn't have this problem since it is transformed
# using the vectorizer fit on training data.
optimal_c, c_auc_train, c_auc_cv = get_optimal_c(X_train_bow_std,
                                                  y_train,
                                                  penalty='l2')
```

```

print("Optimal value of C : {}".format(optimal_c))

train_auc = [(c, train_auc) for c, train_auc in c_auc_train]
cv_auc = [(c, cv_auc) for c, cv_auc in c_auc_cv]

# Plotting c values vs AUC scores

plt.title("C vs AUC")
plt.xlabel("Hyperparameter(c)")
plt.ylabel("Area under ROC curve")
plt.plot(*zip(*train_auc), label='Train_AUC')
plt.plot(*zip(*cv_auc), label='Validation_AUC')
plt.legend()
plt.show()

# Running logistic Regression with optimal alpha value obtained
lr_clf, auc_score, conf_mat = run_lr(X_train_bow_std, y_train,
                                     X_test_bow_std, y_test,
                                     optimal_c, penalty='l2')

print("AUC score:\n {:.2f}".format(auc_score))

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)



auc_score = '%0.2f' % auc_score

# Sparsity of weight vector to results
sparsity = get_sparsity(lr_clf)


# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "BoW", \
                                'l2', optimal_c, sparsity, auc_score]

```

Optimal value of C : 0.001



AUC score:
0.93



```
In [22]: print("Sparsity of weight vector : {}".format(sparsity))
```

Sparsity of weight vector : 0.0%

[5.1.2.1] Performing perturbation test (multicollinearity check) on BOW, SET 1

```
In [180]: # weight vector obtained using l2 regularizer
w = lr_clf.coef_
```

```
# copy of the original train data
X_train_copy = X_train_bow_std.copy()

# Generating random noise to add to X_train
noise = np.random.normal(0,0.0001,
                          X_train_copy.size)
X_train_copy.data += noise
```

```
In [181]: # Running logistic Regression with optimal alpha value obtained
clf_noisy = LogisticRegression(C=optimal_c)
clf_noisy.fit(X_train_copy, y_train)
w_new = clf_noisy.coef_
```



```

In [182]: # There is no non zero element in weight vector
          # so not adding anything to weight vector

          # percentage change in weight vector
          #  $(|W-W'| / |W|) * 100$ 

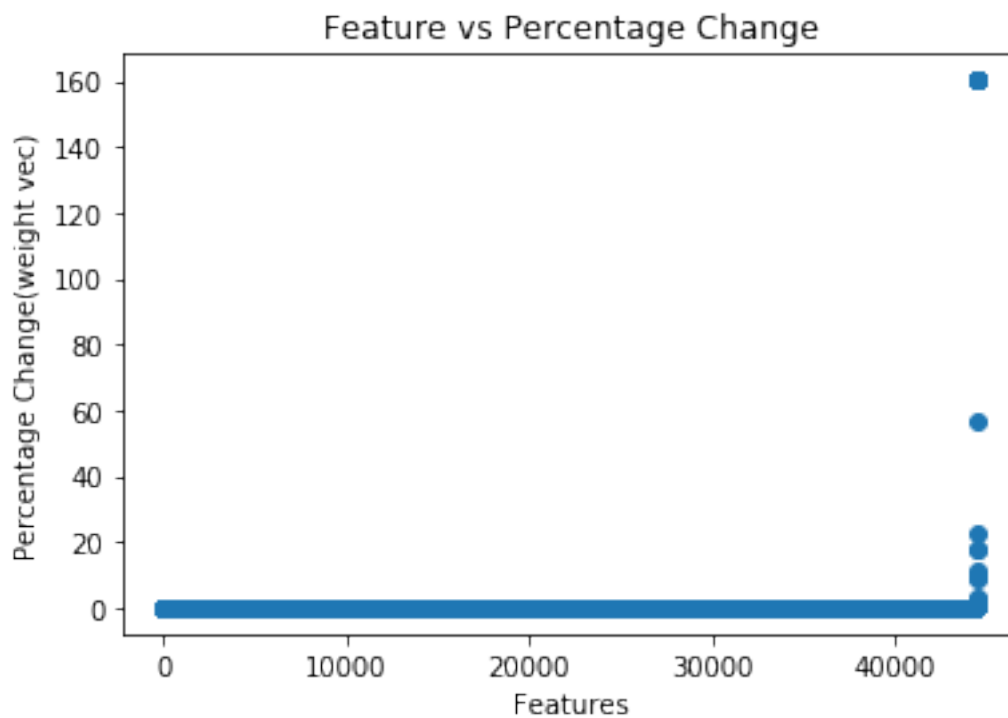
          change_vec = abs((w - w_new)/w)*100

In [183]: change_vec_rounded = np.around(change_vec, decimals=1)
          change_vec_rounded.sort()

In [184]: index = np.arange(0, X_train_copy.shape[1]).reshape(1, X_train_copy.shape[1])

In [185]: plt.scatter(index, change_vec_rounded)
          plt.xlabel("Features")
          plt.ylabel("Percentage Change(weight vec)")
          plt.title('Feature vs Percentage Change')
          plt.show()

```



```

In [186]: for i in range(0,101, 10):
          print("{}th percentile : {}".format(i,np.percentile(change_vec_rounded, i)))

0th percentile : 0.0
10th percentile : 0.0

```

```

20th percentile : 0.0
30th percentile : 0.0
40th percentile : 0.0
50th percentile : 0.0
60th percentile : 0.0
70th percentile : 0.0
80th percentile : 0.0
90th percentile : 0.0
100th percentile : 160.3

```

```

In [192]: i = 99.0
          while(i <= 100.0):
              print("{:.1f}th percentile : {}".format(i,np.percentile(change_vec_rounded, i)))
              i += 0.1

```

```

99.0th percentile : 0.0
99.1th percentile : 0.1
99.2th percentile : 0.1
99.3th percentile : 0.1
99.4th percentile : 0.1
99.5th percentile : 0.1
99.6th percentile : 0.1
99.7th percentile : 0.2
99.8th percentile : 0.2
99.9th percentile : 0.5
100.0th percentile : 160.3

```

```

In [204]: # There is a sudden change in weight vector
          # after the 99.9th percentile to 100th percentile

```

```

          indices = np.where(change_vec > 0.5)[1]

```

```

In [206]: # Loading the saved count vectorizer object
          count_vec = pickle.load(open("count_vec.pkl", 'rb'))

          # Words where change was greater than 0.5
          print(np.take(count_vec.get_feature_names(), indices))

```

```

['amount good' 'becom one' 'bit flavor' 'blurb green' 'border colli'
'cafein darker' 'coffe top' 'color subtl' 'cream flavor' 'dont even'
'enjoy milk' 'entir day' 'envelop help' 'favorit darjeel' 'fresh period'
'general green' 'gob' 'grain' 'like great' 'loaf bread' 'love indian'
'mayb first' 'mex' 'mix well' 'natur oil' 'note one' 'numi love'
'offer twine' 'order whole' 'pipe hot' 'pound' 'prepar food'
'product blurb' 'product process' 'say enjoy' 'settl two' 'shake not'
'sweet tea' 'tast chewi' 'tex mex' 'use give' 'would take']

```

6.1.3 [5.1.3] Feature Importance on BOW, SET 1

```
In [207]: # Loading the saved count vectorizer object
count_vec = pickle.load(open("count_vec.pkl", 'rb'))
```

[5.1.3.1] Top 10 important features of positive class from SET 1

```
In [277]: # Top 10 features in positive class
max_ind_pos = np.argsort(-lr_clf.coef_)[0,:10]
print(np.take(count_vec.get_feature_names(), max_ind_pos))

['love' 'great' 'best' 'delici' 'good' 'perfect' 'excel' 'favorit' 'tasti'
 'high recommend']
```

[5.1.3.2] Top 10 important features of negative class from SET 1

```
In [278]: # Top 10 features in negative class
max_ind_neg = (np.argsort(lr_clf.coef_))[0,:10]
print(np.take(count_vec.get_feature_names(), max_ind_neg))

['disappoint' 'not buy' 'not recommend' 'not good' 'worst' 'not worth'
 'terribl' 'aw' 'horribl' 'wast money']
```

6.2 [5.2] Logistic Regression on TFIDF, SET 2

6.2.1 [5.2.1] Applying Logistic Regression with L1 regularization on TFIDF, SET 2

```
In [28]: # Load the saved vectorized data for train-test datapoints
X_train_tfidf = pickle.load(open('train_tfidf.pkl', 'rb'))
X_test_tfidf = pickle.load(open('test_tfidf.pkl', 'rb'))

std = StandardScaler(with_mean=False)

# Standardizing the vectors
X_train_tfidf_std = std.fit_transform(X_train_tfidf)
X_test_tfidf_std = std.transform(X_test_tfidf)

# Getting an optimal value of hyperparameter c and AUC scores
# This data is used to plot a graph of C-values vs AUC
# There will be a problem of data leakage while using
# Gridsearch on train data but no way to get around it.
# Test data doesn't have this problem since it is transformed
# using the vectorizer fit on training data.
optimal_c, c_auc_train, c_auc_cv = get_optimal_c(X_train_tfidf_std,
                                                  y_train,
                                                  penalty='l1')
```

```

print("Optimal value of C : {}".format(optimal_c))

train_auc = [(c, train_auc) for c, train_auc in c_auc_train]
cv_auc = [(c, cv_auc) for c, cv_auc in c_auc_cv]

# Plotting c values vs AUC scores

plt.title("C vs AUC")
plt.xlabel("Hyperparameter(c)")
plt.ylabel("Area under ROC curve")
plt.plot(*zip(*train_auc), label='Train_AUC')
plt.plot(*zip(*cv_auc), label='Validation_AUC')
plt.legend()
plt.show()

# Running logistic Regression with optimal alpha value obtained
lr_clf, auc_score, conf_mat = run_lr(X_train_tfidf_std, y_train,
                                     X_test_tfidf_std, y_test,
                                     optimal_c, penalty='l1')

print("AUC score:\n {:.2f}".format(auc_score))

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)


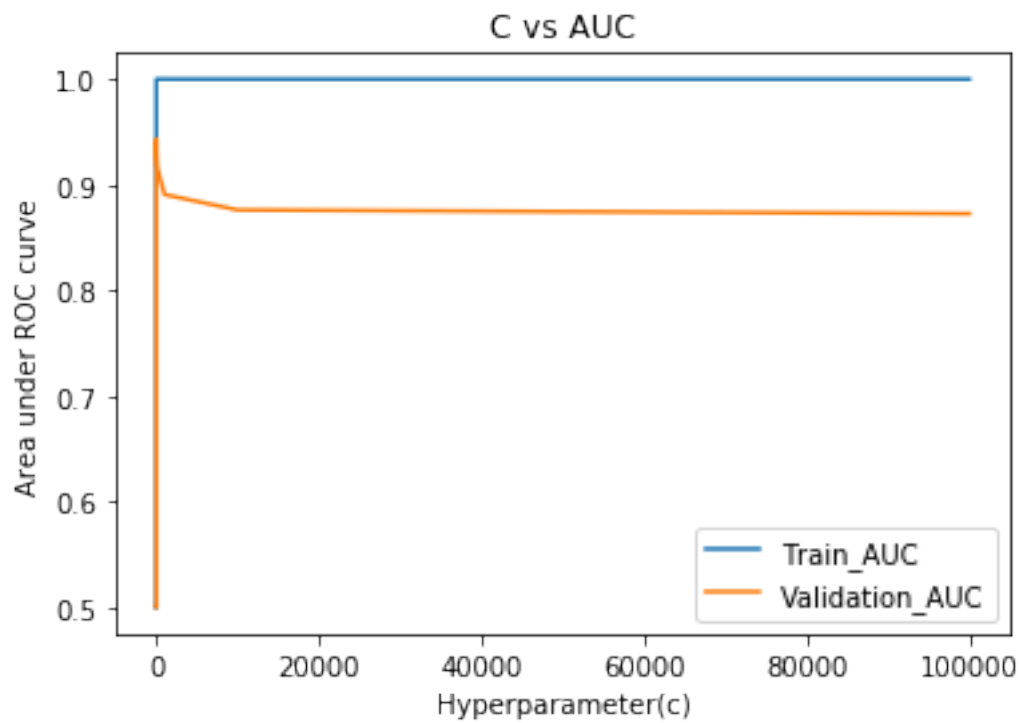
auc_score = '%0.2f' % auc_score

# Sparsity of weight vector to results
sparsity = get_sparsity(lr_clf)


# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "Tf-Idf", \
                                'l1', optimal_c, sparsity, auc_score]

```

Optimal value of C : 0.01



AUC score:
0.96



6.2.2 [5.2.2] Applying Logistic Regression with L2 regularization on TFIDF, SET 2

In [29]: *# Changing the regularizer to l2*

```
# Getting an optimal value of hyperparameter c and AUC scores
# This data is used to plot a graph of C-values vs AUC
# There will be a problem of data leakage while using
# Gridsearch on train data but no way to get around it.
# Test data doesn't have this problem since it is transformed
# using the vectorizer fit on training data.
optimal_c, c_auc_train, c_auc_cv = get_optimal_c(X_train_tfidf_std,
                                                y_train,
                                                penalty='l2')

print("Optimal value of C : {}".format(optimal_c))

train_auc = [(c, train_auc) for c, train_auc in c_auc_train]
cv_auc = [(c, cv_auc) for c, cv_auc in c_auc_cv]

# Plotting c values vs AUC scores
```

```

plt.title("C vs AUC")
plt.xlabel("Hyperparameter(c)")
plt.ylabel("Area under ROC curve")
plt.plot(*(zip(*train_auc)), label='Train_AUC')
plt.plot(*(zip(*cv_auc)), label='Validation_AUC')
plt.legend()
plt.show()

# Running logistic Regression with optimal alpha value obtained
lr_clf, auc_score, conf_mat = run_lr(X_train_tfidf_std, y_train,
                                     X_test_tfidf_std, y_test,
                                     optimal_c, penalty='l2')

print("AUC score:\n {:.2f}".format(auc_score))

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)



auc_score = '%0.2f' % auc_score

# Sparsity of weight vector to results
sparsity = get_sparsity(lr_clf)


# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "Tf-Idf", \
                                'l2', optimal_c, sparsity, auc_score]

```

Optimal value of C : 0.0001



AUC score:
0.95



6.2.3 [5.2.3] Feature Importance on TFIDF, SET 2

```
In [284]: # Loading the saved tf-idf vectorizer object
          tf_idf = pickle.load(open("tf_idf.pkl", 'rb'))
```

[5.2.3.1] Top 10 important features of positive class from SET 2

```
In [285]: # Top 10 features in positive class
          max_ind_pos = np.argsort(-lr_clf.coef_)[0,:10]
          print(np.take(count_vec.get_feature_names(), max_ind_pos))
```

```
['love' 'great' 'best' 'good' 'delici' 'perfect' 'excel' 'favorit' 'use'
 'enjoy']
```

[5.2.3.2] Top 10 important features of negative class from SET 2

```
In [286]: # Top 10 features in negative class
          max_ind_neg = (np.argsort(lr_clf.coef_))[0,:10]
          print(np.take(count_vec.get_feature_names(), max_ind_neg))
```

```
['disappoint' 'not buy' 'not recommend' 'worst' 'not good' 'aw'  
'not worth' 'terribl' 'horribl' 'wast money']
```

6.3 [5.3] Logistic Regression on AVG W2V, SET 3

6.3.1 [5.3.1] Applying Logistic Regression with L1 regularization on AVG W2V SET 3

```
In [30]: # Load the saved vectorized data for train-test datapoints  
X_train_avgw2v = pickle.load(open('train_avgw2v.pkl', 'rb'))  
X_test_avgw2v = pickle.load(open('test_avgw2v.pkl', 'rb'))  
  
std = StandardScaler()  
  
# Standardizing the vectors  
X_train_avgw2v_std = std.fit_transform(X_train_avgw2v)  
X_test_avgw2v_std = std.transform(X_test_avgw2v)  
  
# Getting an optimal value of hyperparameter c and AUC scores  
# This data is used to plot a graph of C-values vs AUC  
# There will be a problem of data leakage while using  
# Gridsearch on train data but no way to get around it.  
# Test data doesn't have this problem since it is transformed  
# using the vectorizer fit on training data.  
optimal_c, c_auc_train, c_auc_cv = get_optimal_c(X_train_avgw2v_std,  
                                                y_train,  
                                                penalty='l1')  
  
print("Optimal value of C : {}".format(optimal_c))  
  
train_auc = [(c, train_auc) for c, train_auc in c_auc_train]  
cv_auc = [(c, cv_auc) for c, cv_auc in c_auc_cv]  
  
# Plotting c values vs AUC scores  
  
plt.title("C vs AUC")  
plt.xlabel("Hyperparameter(c)")  
plt.ylabel("Area under ROC curve")  
plt.plot(*zip(*train_auc), label='Train_AUC')  
plt.plot(*zip(*cv_auc), label='Validation_AUC')  
plt.legend()  
plt.show()  
  
# Running logistic Regression with optimal alpha value obtained  
lr_clf, auc_score, conf_mat = run_lr(X_train_avgw2v_std, y_train,  
                                     X_test_avgw2v_std, y_test,  
                                     optimal_c, penalty='l1')  
  
print("AUC score:\n {:.2f}".format(auc_score))
```

```

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)


auc_score = '%0.2f' % auc_score


# Sparsity of weight vector to results
sparsity = get_sparsity(lr_clf)

# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "Average W2V", \
                                'l1', optimal_c, sparsity, auc_score]


```

Optimal value of C : 1





AUC score:
0.92



6.3.2 [5.3.2] Applying Logistic Regression with L2 regularization on AVG W2V, SET 3

In [31]: *# Changing the regularizer to l2*

```
# Getting an optimal value of hyperparameter c and AUC scores
# This data is used to plot a graph of C-values vs AUC
# There will be a problem of data leakage while using
# Gridsearch on train data but no way to get around it.
# Test data doesn't have this problem since it is transformed
# using the vectorizer fit on training data.
optimal_c, c_auc_train, c_auc_cv = get_optimal_c(X_train_avgw2v_std,
                                                y_train,
                                                penalty='l2')

print("Optimal value of C : {}".format(optimal_c))

train_auc = [(c, train_auc) for c, train_auc in c_auc_train]
cv_auc = [(c, cv_auc) for c, cv_auc in c_auc_cv]

# Plotting c values vs AUC scores

plt.title("C vs AUC")
plt.xlabel("Hyperparameter(c)")
plt.ylabel("Area under ROC curve")
plt.plot(*zip(*train_auc), label='Train_AUC')
plt.plot(*zip(*cv_auc), label='Validation_AUC')
plt.legend()
plt.show()

# Running logistic Regression with optimal alpha value obtained
lr_clf, auc_score, conf_mat = run_lr(X_train_avgw2v_std, y_train,
                                     X_test_avgw2v_std, y_test,
                                     optimal_c, penalty='l2')

print("AUC score:\n {:.2f}".format(auc_score))

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)


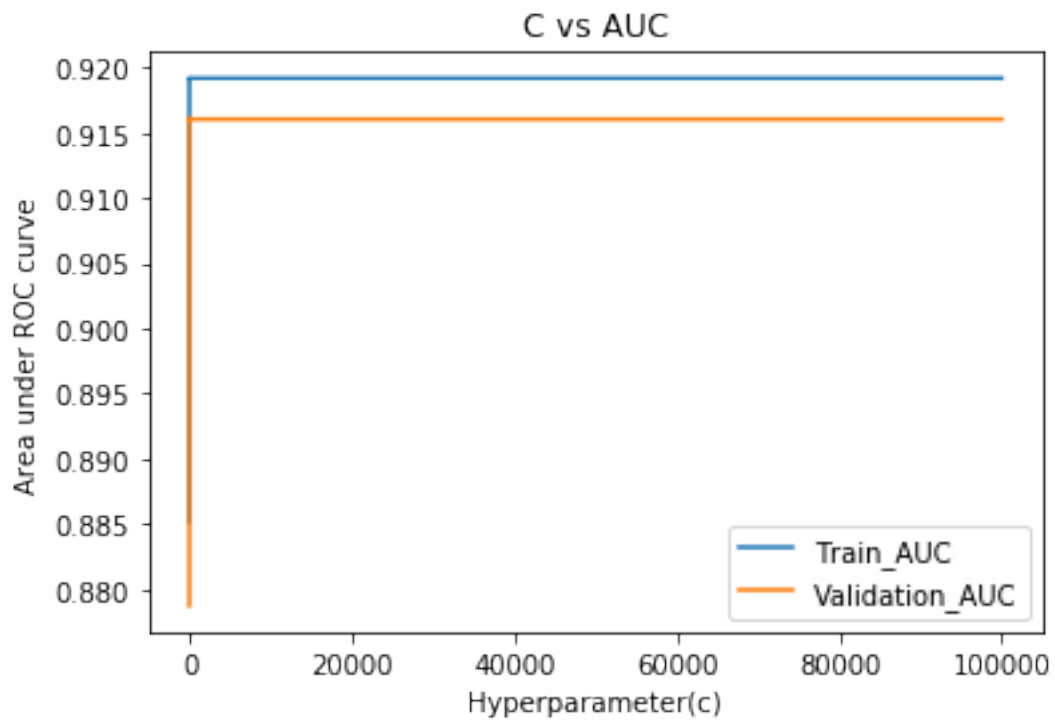
auc_score = '%0.2f' % auc_score

# Sparsity of weight vector to results
sparsity = get_sparsity(lr_clf)


# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "Average W2V", \
```

```
'12', optimal_c, sparsity, auc_score]
```

Optimal value of C : 0.1



AUC score:
0.92



6.4 [5.4] Logistic Regression on TFIDF W2V, SET 4

6.4.1 [5.4.1] Applying Logistic Regression with L1 regularization on TFIDF W2V, SET 4

```
In [32]: # Load the saved vectorized data for train-test datapoints
X_train_tfidf2v = pickle.load(open('train_tfidf2v.pkl', 'rb'))
X_test_tfidf2v = pickle.load(open('test_tfidf2v.pkl', 'rb'))

std = StandardScaler()

# Standardizing the vectors
X_train_tfidf2v_std = std.fit_transform(X_train_tfidf2v)
X_test_tfidf2v_std = std.transform(X_test_tfidf2v)

# Getting an optimal value of hyperparameter c and AUC scores
# This data is used to plot a graph of C-values vs AUC
```

```

# There will be a problem of data leakage while using
# Gridsearch on train data but no way to get around it.
# Test data doesn't have this problem since it is transformed
# using the vectorizer fit on training data.
optimal_c, c_auc_train, c_auc_cv = get_optimal_c(X_train_tfidfv2v_std,
                                                y_train,
                                                penalty='l1')

print("Optimal value of C : {}".format(optimal_c))

train_auc = [(c, train_auc) for c, train_auc in c_auc_train]
cv_auc = [(c, cv_auc) for c, cv_auc in c_auc_cv]

# Plotting c values vs AUC scores

plt.title("C vs AUC")
plt.xlabel("Hyperparameter(c)")
plt.ylabel("Area under ROC curve")
plt.plot(*zip(*train_auc), label='Train_AUC')
plt.plot(*zip(*cv_auc), label='Validation_AUC')
plt.legend()
plt.show()

# Running logistic Regression with optimal alpha value obtained
lr_clf, auc_score, conf_mat = run_lr(X_train_tfidfv2v_std, y_train,
                                     X_test_tfidfv2v_std, y_test,
                                     optimal_c, penalty='l1')

print("AUC score:\n {:.2f}".format(auc_score))

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)


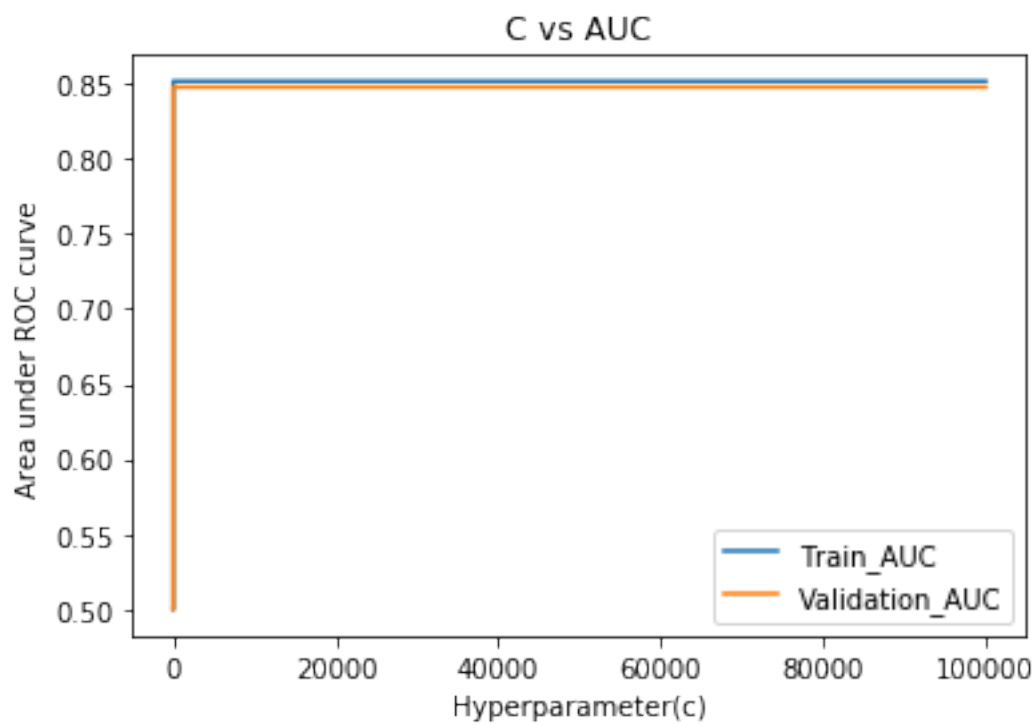
auc_score = '%0.2f' % auc_score

# Sparsity of weight vector to results
sparsity = get_sparsity(lr_clf)


# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "Tfidf_W2V", \
                                'l1', optimal_c, sparsity, auc_score]

```

Optimal value of C : 1



AUC score:
0.85



6.4.2 [5.4.2] Applying Logistic Regression with L2 regularization on TFIDF W2V, SET 4

In [33]: *# Changing the regularizer to l2*

```
# Getting an optimal value of hyperparameter c and AUC scores
# This data is used to plot a graph of C-values vs AUC
# There will be a problem of data leakage while using
# Gridsearch on train data but no way to get around it.
# Test data doesn't have this problem since it is transformed
# using the vectorizer fit on training data.
optimal_c, c_auc_train, c_auc_cv = get_optimal_c(X_train_tfidf_w2v_std,
                                                y_train,
                                                penalty='l2')
print("Optimal value of C : {}".format(optimal_c))

train_auc = [(c, train_auc) for c, train_auc in c_auc_train]
cv_auc = [(c, cv_auc) for c, cv_auc in c_auc_cv]

# Plotting c values vs AUC scores
```

```

plt.title("C vs AUC")
plt.xlabel("Hyperparameter(c)")
plt.ylabel("Area under ROC curve")
plt.plot(*(zip(*train_auc)), label='Train_AUC')
plt.plot(*(zip(*cv_auc)), label='Validation_AUC')
plt.legend()
plt.show()

# Running logistic Regression with optimal alpha value obtained
lr_clf, auc_score, conf_mat = run_lr(X_train_tfidfw2v_std, y_train,
                                     X_test_tfidfw2v_std, y_test,
                                     optimal_c, penalty='l2')

print("AUC score:\n {:.2f}".format(auc_score))

# Plotting confusion matrix
plot_confusion_matrix(conf_mat)



auc_score = '%0.2f' % auc_score

# Sparsity of weight vector to results
sparsity = get_sparsity(lr_clf)


# Adding the results to our results dataframe
results.loc[results.shape[0]] = ["Review Text", "TfIdf_W2V", \
                                'l2', optimal_c, sparsity, auc_score]

```

Optimal value of C : 1



AUC score:
0.85



7 [6] Conclusions

1. We tried BoW, TF-IDF, Average Word2Vec and Tfidf weighted Word2Vec vectorizers on Logistic Regression using L1 and L2 regularizer.
2. AUC score for Logistic Regression was better using L1 regularization for BoW and TFIDF vectorizers.
3. Sparsity in weight vector using BoW vectorizer with l1 regularization --> 91.40%
4. Sparsity in weight vector using TFIDF vectorizer with l1 regularization --> 90.56%
5. BoW : Top 10 important features of positive class 'love' 'great' 'best' 'delici' 'good' 'perfect' 'excel' 'favorit' 'tasti' 'high recommend'
Top 10 important features of negative class 'disappoint' 'not buy' 'not recommend' 'not good' 'worst' 'not worth' 'terribl' 'aw' 'horribl' 'wast money'
6. TFIDF : Top 10 important features of positive class 'love' 'great' 'best' 'good' 'delici' 'perfect' 'excel' 'favorit' 'use' 'enjoy'
Top 10 important features of negative class 'disappoint' 'not buy' 'not recommend' 'worst' 'not good' 'aw' 'not worth' 'terribl' 'horribl' 'wast money'

7. Pertubation test : Observed a minimal change in weight vectors after adding noise to training data

99.0th percentile : 0.0

99.1th percentile : 0.1

99.2th percentile : 0.1

99.3th percentile : 0.1

99.4th percentile : 0.1

99.5th percentile : 0.1

99.6th percentile : 0.1

99.7th percentile : 0.2

99.8th percentile : 0.2

99.9th percentile : 0.5

100.0th percentile : 160.3

Words for which change was more than 0.5% (99.9th percentile)

'amount good' 'becom one' 'bit flavor' 'blurb green' 'border colli' 'caffein darker' 'coffe top'
'color subtl' 'cream flavor' 'dont even' 'enjoy milk' 'entir day' 'envelop help' 'favorit dar-
jeel' 'fresh period' 'general green' 'gob' 'grain' 'like great' 'loaf bread' 'love indian' 'mayb
first' 'mex' 'mix well' 'natur oil' 'note one' 'numi love' 'offer twine' 'order whole' 'pipe hot'
'pound' 'prepar food' 'product blurb' 'product process' 'say enjoy' 'settl two' 'shake not'
'sweet tea' 'tast chewi' 'tex mex' 'use give' 'would take'

In [34]: print_results(results)

Features-Used	Vectorizer	Regularizer	C	Sparsity	AUC
Review Text	BoW	11	0.01	91.38	0.95
Review Text	BoW	12	0.001	0.0	0.93
Review Text	Tf-Idf	11	0.01	90.59	0.96
Review Text	Tf-Idf	12	0.0001	0.0	0.95
Review Text	Average W2V	11	1.0	0.0	0.92
Review Text	Average W2V	12	0.1	0.0	0.92
Review Text	TfIdf_W2V	11	1.0	0.0	0.85
Review Text	TfIdf_W2V	12	1.0	0.0	0.85